

Xunmei (May) Liu

Seattle, WA | 206-419-3458 | liuxmmay2021@outlook.com | [LinkedIn](#)

Education

University of Washington

M.S. in Computer Science, GPA: 4.0/4.0

Seattle, WA

Expected Sep. 2025 – Dec. 2026

University of Washington

B.S. in Computer Science, GPA: 3.90/4.0

Seattle, WA

Sep. 2021 – June 2025

Relevant Courses: Deep Learning, NLP, Robotics, AI, Machine Learning, Distributed Systems, Operating Systems

Teaching: Teacher Assistant for Probability, Data Structures, Computer Graphics, Intro to Java III for nine quarters

Skills

Programming: Python, Java, PyTorch, Go, C/C++, JavaScript, SQL

Systems & Tools: Redis, React, Ray, Kubernetes, FastAPI, RocksDB, HDFS

ML Areas: RLHF, large language models, distributed training, vision-language models, computer vision, recommendation systems, retrieval-augmented generation, agents

Work Experience

Google

Software Engineer Intern, Spanner

Sunnyvale, CA

Expected June 2026 – Aug. 2026

- Train supervised learning models to optimize CPU provisioning for replication workloads within Spanner's infrastructure.

ByteDance Seed

Software Engineer Intern, Seed-Infra-Training

Seattle, WA

Feb. 2025 – Aug. 2025

- Developed and deployed a high-performance replay buffer in Seed's internal verl, resolving memory bottlenecks in large-scale rollout pipelines.
- Implemented CRUD APIs and contributed to the architecture design of a code and terminal agent sandbox.
- Built asynchronous push/get/sample/delete APIs using a hybrid LRU in-memory cache and RocksDB.
- Supported customizable sampling strategies with efficient in-memory indexing and periodic HDFS backups.

AdasEco

Machine Learning Intern

Remote

June 2024 – Aug. 2024

- Trained a YOLOv8 model for autonomous driving object detection in PyTorch and achieved a 20% mAP gain with 50% faster training through distributed data-parallel optimization.

Publications

VFIG: Vectorizing Complex Figures in SVG with Vision-Language Models (ArXiv | Project website) 2026

- Qijia He, **Xunmei Liu**, Hammaad Memon, Ziang Li, Zixian Ma, Jaemin Cho, Jason Ren, Daniel S. Weld, Ranjay Krishna. *Co-first author*. Submitted to ECCV 2026.
- Main contributions: reinforcement learning training, reward design, data generation support, and benchmark evaluation.

eKichabi v2: Designing and Scaling a Dual-Platform Agricultural Technology in Rural Tanzania 2024

- Ananditha Raghunath, Alexander Metzger, Hans Easton, **Xunmei Liu**, Fanchong Wang, Yunqi Wang, Yunwei Zhao, Hosea Mpogole, Richard Anderson. *Co-second author*.
- Contributed to the system and research effort behind eKichabi v2, including the USSD platform and backend data workflow.

Ongoing Research Experience

University of Washington, RAIVN Lab

Research Assistant, Target: ICLR 2026, First Author

Seattle, WA

Mar. 2026 – Present

- Built a reinforcement learning pipeline to train MolmoWeb-4B as a multimodal web agent, enabling it to interact with live websites using browser screenshots and page content.
- Applied web-agent RL training to real-world WebArena-Infinity tasks across applications such as Gmail, Reddit, and GitLab, targeting multi-step website interactions.